

What Is A Record?

- Recall:
 - Records: If I_1, \dots, I_k are distinct identifiers, and T_1, \dots, T_k are type expressions, then **RECORD**($I_1: T_1, \dots, I_k: T_k$) is a type expression denoting a record type with k named fields.
- In other words:
 - A record is an aggregation of potentially heterogeneous (sub-)objects into a single unit.
 - The aggregated objects are called the fields of the record, and are referenced by name.

What Is A Record?

- Recall:
 - Records: If I_1, \dots, I_k are distinct identifiers, and T_1, \dots, T_k are type expressions, then **RECORD**($I_1: T_1, \dots, I_k: T_k$) is a type expression denoting a record type with k named fields.
- In other words:
 - A record is an aggregation of potentially heterogeneous (sub-)objects into a single unit.
 - The aggregated objects are called the fields of the record, and are referenced by name.
- [C] Record types are called structure types (written `struct`).
 - Can be **tagged** (`struct T {...} x;`) or **untagged** (`struct {...} x;`).
 - Often used with `typedef` (`typedef struct {...} T;`) to produce new type names for structure types.

What Is A Record?

- Recall:
 - Records: If I_1, \dots, I_k are distinct identifiers, and T_1, \dots, T_k are type expressions, then **RECORD**($I_1: T_1, \dots, I_k: T_k$) is a type expression denoting a record type with k named fields.
- In other words:
 - A record is an aggregation of potentially heterogeneous (sub-)objects into a single unit.
 - The aggregated objects are called the fields of the record, and are referenced by name.
- [C] Record types are called structure types (written `struct`).
 - Can be tagged (`struct T {...} x;`) or untagged (`struct {...} x;`).
 - Often used with `typedef` (`typedef struct {...} T;`) to produce new type names for structure types.

```
typedef struct point {  
    int x, y;  
} Point;
```

```
typedef struct {  
    struct point start;  
    Point end;  
} Line;
```

```
Point p, *pptr;  
Line ln[5];
```

Run-Time Representation of Records

- Assume that we have the following two functions available.
 - *size*: $\text{Type} \rightarrow \mathbf{int}$, giving the number of contiguous bytes of memory that a variable of a given type occupies at run-time.
 - *addr*: $\text{Identifier} \rightarrow \mathbf{int}$, giving the starting address in memory of this chunk for a given identifier.

Run-Time Representation of Records

- Assume that we have the following two functions available.
 - $size: \text{Type} \rightarrow \mathbf{int}$, giving the number of contiguous bytes of memory that a variable of a given type occupies at run-time.
 - $addr: \text{Identifier} \rightarrow \mathbf{int}$, giving the starting address in memory of this chunk for a given identifier.
- Then, variable x of type T occupies memory bytes $addr(x)$ through $addr(x) + size(x) - 1$.

Run-Time Representation of Records

- Assume that we have the following two functions available.
 - $size: \text{Type} \rightarrow \mathbf{int}$, giving the number of contiguous bytes of memory that a variable of a given type occupies at run-time.
 - $addr: \text{Identifier} \rightarrow \mathbf{int}$, giving the starting address in memory of this chunk for a given identifier.
- Then, variable x of type T occupies memory bytes $addr(x)$ through $addr(x) + size(x) - 1$.
- If T is a record type, the language definition specifies a mapping of each field f (of type τ_f) of the record type to offsets δ_f , such that $addr(x.f) = addr(x) + \delta_f$.

Run-Time Representation of Records

- Assume that we have the following two functions available.
 - $size: \text{Type} \rightarrow \mathbf{int}$, giving the number of contiguous bytes of memory that a variable of a given type occupies at run-time.
 - $addr: \text{Identifier} \rightarrow \mathbf{int}$, giving the starting address in memory of this chunk for a given identifier.
- Then, variable x of type T occupies memory bytes $addr(x)$ through $addr(x) + size(x) - 1$.
- If T is a record type, the language definition specifies a mapping of each field f (of type τ_f) of the record type to offsets δ_f , such that $addr(x.f) = addr(x) + \delta_f$.
- And we will have $size(T) \geq \sum_f size(\tau_f)$.
 - I.e., each field will be contiguous in memory, but there may be gaps (“padding”) between successive fields.

Operations on Record Types

- We have a single operation defined on a record type: accessing a field.
- Given a variable x of a record type T , what does $x.f$ mean?
 - Depends on context: rvalue or lvalue.

Operations on Record Types

- We have a single operation defined on a record type: accessing a field.
- Given a variable x of a record type T , what does $x.f$ mean?
 - Depends on context: rvalue or lvalue.
- Case 1: Rvalue (i.e., a **reference** $\dots x.f \dots$ within an expression).
 - Returns a value of type τ_f from the bit pattern in memory bytes $addr(x) + \delta(f)$ through $addr(x) + \delta(f) + size(\tau_f) - 1$.

Operations on Record Types

- We have a single operation defined on a record type: accessing a field.
- Given a variable x of a record type T , what does $x.f$ mean?
 - Depends on context: rvalue or lvalue.
- Case 1: Rvalue (i.e., a reference $\dots x.f \dots$ within an expression).
 - Returns a value of type τ_f from the bit pattern in memory bytes $addr(x) + \delta(f)$ through $addr(x) + \delta(f) + size(\tau_f) - 1$.
- Case 2: Lvalue (i.e., an **update** $x.f = \dots$ within an assignment).
 - Memory bytes $addr(x) + \delta(f)$ through $addr(x) + \delta(f) + size(\tau_f) - 1$ with the type-appropriate bit-pattern representation of the rvalue of the assignment statement.

Record Layout and Data Alignment

- Two nagging questions.
 1. Why do we have $size(T) \geq \sum_f size(\tau_f)$ rather than $size(T) = \sum_f size(\tau_f)$?
 2. Why does the C standard insist that “For two structures, corresponding members shall be declared in the same order.” [C17 ballot, §6.2.7]?

Record Layout and Data Alignment

- Two nagging questions.
 1. Why do we have $size(T) \geq \sum_f size(\tau_f)$ rather than $size(T) = \sum_f size(\tau_f)$?
 2. Why does the C standard insist that “For two structures, corresponding members shall be declared in the same order.” [C17 ballot, §6.2.7]?
- Both questions have the same answer: data alignment.

Record Layout and Data Alignment

- Two nagging questions.
 1. Why do we have $size(T) \geq \sum_f size(\tau_f)$ rather than $size(T) = \sum_f size(\tau_f)$?
 2. Why does the C standard insist that “For two structures, corresponding members shall be declared in the same order.” [C17 ballot, §6.2.7]?
- Both questions have the same answer: data alignment.
- What is data alignment?
 - Computers often impose restrictions on the valid/desired placement of data objects in memory (i.e., $addr(x)$) in order to simplify high-performance memory design. These restrictions are called **data alignment rules**.
 - Such restrictions are not particularly onerous, because they can be easily handled by the compiler and the run-time system.
 - Reality check: The x64 architecture will work correctly regardless of data alignment (except for some SSE instructions), *but performance may suffer*.

Data Alignment in C

- A variable x of a basic data type of size K bytes and located at memory address μ_x is said to be aligned iff $\mu_x = 0 \bmod K$.

Data Alignment in C

- A variable x of a basic data type of size K bytes and located at memory address μ_x is said to be aligned iff $\mu_x = 0 \bmod K$.
- A variable x of a derived data type T located at memory address μ_x iff the following two conditions hold:
 - [The sub-object rule] Every sub-object of this type is (recursively) aligned.
 - [The array rule] Every element of an object y of type “array of T ” is aligned.

Data Alignment in C

- A variable x of a basic data type of size K bytes and located at memory address μ_x is said to be aligned iff $\mu_x = 0 \bmod K$.
- A variable x of a derived data type T located at memory address μ_x iff the following two conditions hold:
 - [The sub-object rule] Every sub-object of this type is (recursively) aligned.
 - [The array rule] Every element of an object y of type “array of T ” is aligned.

```
typedef struct point {  
    int x, y;  
} Point;
```

```
typedef struct {  
    struct point start;  
    Point end;  
} Line;
```

```
Line ln[5];
```


Data Alignment in C

- A variable x of a basic data type of size K bytes and located at memory address μ_x is said to be aligned iff $\mu_x = 0 \bmod K$.
- A variable x of a derived data type T located at memory address μ_x iff the following two conditions hold:
 - [The sub-object rule] Every sub-object of this type is (recursively) aligned.
 - [The array rule] Every element of an object y of type “array of T ” is aligned.

```
struct S1 {  
    int i;  
    char c;  
    int j;  
};
```

```
struct S2 {  
    int i;  
    int j;  
    char c;  
};
```

```
typedef struct point {  
    int x, y;  
} Point;
```

```
typedef struct {  
    struct point start;  
    Point end;  
} Line;
```

```
Line ln[5];
```

Function-Valued Fields in Records

- Nothing prevents us from adding function pointers as fields in records.

Function-Valued Fields in Records

- Nothing prevents us from adding function pointers as fields in records.

```
typedef struct point {  
    int x, y;  
    double (*d)(double, double);  
} Point;
```

```
typedef struct {  
    struct point start;  
    Point end;  
    double (*l)(Point, Point);  
} Line;
```

Function-Valued Fields in Records

- Nothing prevents us from adding function pointers as fields in records.

```
typedef struct point {          #include <math.h>
    int x, y;                  double L2_dist(double x, double y) {
    double (*d)(double, double); return sqrt(x*x+y*y);
} Point;                        }
                                double L1_dist(double x, double y) {
                                return fabs(x) + fabs(y);
typedef struct {                }
    struct point start;        double line_len(Point p, Point q) {
    Point end;                  return L2_dist(L1_dist(p.x, p.y),
    double (*l)(Point, Point);  L1_dist(q.x, q.y));
} Line;                          }
```

Function-Valued Fields in Records

- Nothing prevents us from adding function pointers as fields in records.

```
typedef struct point {      #include <math.h>
    int x, y;              double L2_dist(double x, double y) {
    double (*d)(double, double); return sqrt(x*x+y*y);
} Point;                  }
                           double L1_dist(double x, double y) {
                           return fabs(x) + fabs(y);
typedef struct {           }
    struct point start;    }
    Point end;             double line_len(Point p, Point q) {
    double (*l)(Point, Point); return L2_dist(L1_dist(p.x, p.y),
} Line;                   L1_dist(q.x, q.y));
                           }

    Point p1, p2;
    Line l;
    p1.x = 10; p1.y = 20; p1.d = L2_dist;
    p2.x = 30; p2.y = 40; p2.d = L1_dist;
    l.start = p1; l.end = p2; l.len = line_len;
```

Function-Valued Fields in Records

- Nothing prevents us from adding function pointers as fields in records.
- What is this beginning to resemble?

```
typedef struct point {      #include <math.h>
    int x, y;               double L2_dist(double x, double y) {
    double (*d)(double, double); return sqrt(x*x+y*y);
} Point;                   }
                           double L1_dist(double x, double y) {
typedef struct {             return fabs(x) + fabs(y);
    struct point start;     }
    Point end;              double line_len(Point p, Point q) {
    double (*l)(Point, Point); return L2_dist(L1_dist(p.x, p.y),
} Line;                     L1_dist(q.x, q.y));
                           }

    Point p1, p2;
    Line l;
    p1.x = 10; p1.y = 20; p1.d = L2_dist;
    p2.x = 30; p2.y = 40; p2.d = L1_dist;
    l.start = p1; l.end = p2; l.len = line_len;
```