

The Design Problem

- Since the only thing the allocator can work with is the collection of free blocks, the management of this free space (aka the **free list**) is the fundamental design problem.

The Design Problem

- Since the only thing the allocator can work with is the collection of free blocks, the management of this free space (aka the free list) is the fundamental design problem.
- Four design dimensions.
 - Free list **organization**: How to keep track of the free blocks.
 - Free block **selection**: How to pick an appropriate free block to satisfy the current allocation request from the mutator.
 - Free block **splitting**: What to do with the remainder of a free block after allocation.
 - Free block **coalescing**: How to handle an allocated block that is being returned to the free list.

The Design Problem

- Since the only thing the allocator can work with is the collection of free blocks, the management of this free space (aka the free list) is the fundamental design problem.
- Four design dimensions.
 - Free list organization: How to keep track of the free blocks.
 - Free block selection: How to pick an appropriate free block to satisfy the current allocation request from the mutator.
 - Free block splitting: What to do with the remainder of a free block after allocation.
 - Free block coalescing: How to handle an allocated block that is being returned to the free list.
- The “free list” is conceptual.
 - Everything is just a byte array, and the allocator overlays a particular interpretation of these bytes.
 - Very low-level. No conventional abstractions to help us out.

Huge Design Space

- Many policy choices for each dimension, interacting in complicated ways.

Huge Design Space

- Many policy choices for each dimension, interacting in complicated ways.
 - Free list **organization**: How to keep track of the free blocks.
 - Implicit, explicit, binned.
 - Singly- or doubly-linked.
 - Null-terminated or circular.
 - Returned blocks maintained in LIFO, address-sorted, or size-sorted order.

Huge Design Space

- Many policy choices for each dimension, interacting in complicated ways.
 - Free list organization: How to keep track of the free blocks.
 - Implicit, explicit, binned.
 - Singly- or doubly-linked.
 - Null-terminated or circular.
 - Returned blocks maintained in LIFO, address-sorted, or size-sorted order.
 - Free block **selection**: How to pick an appropriate free block to satisfy the current allocation request from the mutator.
 - Best-fit, worst-fit, first-fit, next-fit.

Huge Design Space

- Many policy choices for each dimension, interacting in complicated ways.
 - Free list organization: How to keep track of the free blocks.
 - Implicit, explicit, binned.
 - Singly- or doubly-linked.
 - Null-terminated or circular.
 - Returned blocks maintained in LIFO, address-sorted, or size-sorted order.
 - Free block selection: How to pick an appropriate free block to satisfy the current allocation request from the mutator.
 - Best-fit, worst-fit, first-fit, next-fit.
 - Free block **splitting**: What to do with the remainder of a free block after allocation.
 - No split, always split, threshold split.
 - Allocate block from bottom or top of free block.

Huge Design Space

- Many policy choices for each dimension, interacting in complicated ways.
 - Free list organization: How to keep track of the free blocks.
 - Implicit, explicit, binned.
 - Singly- or doubly-linked.
 - Null-terminated or circular.
 - Returned blocks maintained in LIFO, address-sorted, or size-sorted order.
 - Free block selection: How to pick an appropriate free block to satisfy the current allocation request from the mutator.
 - Best-fit, worst-fit, first-fit, next-fit.
 - Free block splitting: What to do with the remainder of a free block after allocation.
 - No split, always split, threshold split.
 - Allocate block from bottom or top of free block.
 - Free block **coalescing**: How to handle an allocated block that is being returned to the free list.
 - Immediate, deferred, none.

The K&R Allocator

- From K&R2e, §8.7.
 - Short, portable, performant.
- Very few assumptions.
 - Does not assume that successive calls to `sbrk()` return contiguous memory in process's address space.
 - Only non-portability is the assumption that pointers to different blocks returned by `sbrk ()` can be meaningfully compared.

The K&R Allocator: Design Choices

- Free list **organization**: How to keep track of the free blocks.
 - Implicit, explicit, binned.
 - Singly- or doubly-linked.
 - Null-terminated or circular.
 - Returned blocks maintained in LIFO, address-sorted, or size-sorted order.
- Free block **selection**: How to pick an appropriate free block to satisfy the current allocation request from the mutator.
 - Best-fit, worst-fit, first-fit, next-fit.
- Free block **splitting**: What to do with the remainder of a free block after allocation.
 - No split, always split, threshold split.
 - Allocate block from bottom or top of free block.
- Free block **coalescing**: How to handle an allocated block that is being returned to the free list.
 - Immediate, deferred, none.

The K&R Allocator: Data Structures

```
typedef long Align; /* for alignment to long boundary */

union header { /* block header: */
    struct {
        union header *ptr; /* next block if on free list */
        unsigned size; /* size of this block */
    } s;
    Align x; /* force alignment of blocks */
};

typedef union header Header;

static Header base; /* empty list to get started */
static Header *freep = NULL; /* start of free list */
```

The K&R Allocator: malloc ()

```
/* malloc: general-purpose storage allocator */
void *malloc(unsigned nbytes) {
    Header *p, *prevp;
    Header *morecore(unsigned);
    unsigned nunits;

    nunits = (nbytes+sizeof(Header)-1)/sizeof(Header) + 1;
    if ((prevp = freep) == NULL) { /* no free list yet */
        base.s.ptr = freep = prevp = & base;
        base.s.size = 0;
    }
    for (p = prevp->s.ptr; ; prevp = p, p = p->s.ptr) {
        if (p->s.size >= nunits) { /* big enough */
            if (p->s.size == nunits) /* exactly */
                prevp->s.ptr = p->s.ptr;
            else { /* allocate tail end */
                p->s.size -= nunits;
                p += p->s.size;
                p->s.size = nunits;
            }
            freep = prevp;
            return (void *) (p+1);
        }
    }
    if (p == freep) /* wrapped around free list */
        if ((p = morecore(nunits)) == NULL)
            return NULL; /* none left */
}
```

The K&R Allocator: morecore ()

```
#define NALLOC 1024 /* minimum #units to request */

/* morecore: ask system for more memory */
static Header *morecore(unsigned nu) {
    char *cp, *sbrk(int);
    Header *up;

    if (nu < NALLOC)
        nu = NALLOC;
    cp = sbrk(nu*sizeof(Header));
    if (cp == (char *) -1) /* no space at all */
        return NULL;
    up = (Header *) cp;
    up->s.size = nu;
    free((void *) (up+1));
    return freep;
}
```

The K&R Allocator: free ()

```
/* free: put block ap in free list */
void free(void *ap) {
    Header *bp, *p;

    bp = (Header *)ap - 1; /* point to block header */
    for (p = freep; !(bp > p && bp < p->s.ptr); p = p->s.ptr)
        if (p >= p->s.ptr && (bp > p || bp < p->s.ptr))
            break; /* freed block at start or end of arena */

    if (bp + bp->s.size == p->s.ptr) { /* join to upper nbr */
        bp->s.size += p->s.ptr->s.size;
        bp->s.ptr = p->s.ptr->s.ptr;
    } else
        bp->s.ptr = p->s.ptr;
    if (p + p->s.size == bp) { /* join to lower nbr */
        p->s.size += bp->s.size;
        p->s.ptr = bp->s.ptr;
    } else
        p->s.ptr = bp;
    freep = p;
}
```