

Garbage Collection

- Garbage collection is the automatic reclamation of computer storage.
 - Program does not need to invoke `free()` or `delete()` methods to explicitly reclaim heap memory.

Garbage Collection

- Garbage collection is the automatic reclamation of computer storage.
 - Program does not need to invoke `free()` or `delete()` methods to explicitly reclaim heap memory.
- The garbage collector's role is to identify data objects that are no longer in use and make their space available for reuse by the mutator.


Garbage Collection

- Garbage collection is the automatic reclamation of computer storage.
 - Program does not need to invoke `free()` or `delete()` methods to explicitly reclaim heap memory.
- The garbage collector's role is to identify data objects that are no longer in use and make their space available for reuse by the mutator.
- An object is considered to be **garbage** and subject to reclamation if it is not reachable by the mutator via any path of pointer traversals.
 - Objects that are potentially reachable through such paths of pointer traversals are said to be **live**.
 - Liveness is a global property of an executing process.
 - Note that this run-time notion of liveness is very different from the compile-time notion of liveness that we discussed in the context of register allocation.

Motivation for Garbage Collection


- Necessary for fully modular programming.
 - A method operating on a data structure should by default not need to know what other methods may be operating on the same data structure.
 - If explicit deallocation is the norm, then some module must be responsible for knowing when other modules are no longer interested in the object – thereby introducing unnecessary cross-module dependencies.
- Unnecessary complications of explicit memory management can:
 - Break the basic abstractions of the programming language.
 - Lead to slow memory leaks.
 - Create “heisenbugs”.
 - Make code brittle.

Garbage Collection In The Block Diagram




Garbage Collection In The Block Diagram

- The run-time allocation routines perform special actions to reclaim space, as necessary, when an allocation request is not easily satisfied.
 - Calls to the “deallocator” are implicit in calls to the allocator.



Garbage Collection In The Block Diagram

- The run-time allocation routines perform special actions to reclaim space, as necessary, when an allocation request is not easily satisfied.
 - Calls to the “deallocator” are implicit in calls to the allocator.
- Most collectors require some cooperation from the compiler or interpreter.
 - Object formats must be recognizable by the garbage collector.
 - Certain invariants must be preserved by the running code.
 - These requirements do not prevent the use of compiler optimizations.



The Two-Phase Abstraction

- Conceptually, the functioning of a garbage collector consists of two distinct phases (i.e., sets of activities).
 - **Garbage detection**, i.e., distinguishing live objects from garbage objects through some means.
 - **Storage reclamation**, i.e., reclaiming the storage of the garbage objects so that it is available for use by the mutator.

The Two-Phase Abstraction

- Conceptually, the functioning of a garbage collector consists of two distinct phases (i.e., sets of activities).
 - Garbage detection, i.e., distinguishing live objects from garbage objects through some means.
 - Storage reclamation, i.e., reclaiming the storage of the garbage objects so that it is available for use by the mutator.
- In practice, these two phases may be functionally or temporally interleaved.

Conservatism in Garbage Collectors

- In order to declare an object as being garbage, the collector needs to be able to prove that it will never be used again.
 - An ideal garbage collector would be able to reclaim every object's space just after the last use of the object.

Conservatism in Garbage Collectors

- In order to declare an object as being garbage, the collector needs to be able to prove that it will never be used again.
 - An ideal garbage collector would be able to reclaim every object's space just after the last use of the object.
- In general, the problem of determining exactly which objects are live is undecidable.

Conservatism in Garbage Collectors

- In order to declare an object as being garbage, the collector needs to be able to prove that it will never be used again.
 - An ideal garbage collector would be able to reclaim every object's space just after the last use of the object.
- In general, the problem of determining exactly which objects are live is undecidable.
- All garbage collectors use *some efficient but conservative approximation to liveness*.
 - In reference counting, the approximation is that an object can't be live unless it is referenced.
 - In tracing garbage collection, the approximation is that an object can't be live unless it is reachable.

Conservatism in Garbage Collectors

- In order to declare an object as being garbage, the collector needs to be able to prove that it will never be used again.
 - An ideal garbage collector would be able to reclaim every object's space just after the last use of the object.
- In general, the problem of determining exactly which objects are live is undecidable.
- All garbage collectors use *some efficient but conservative approximation to liveness*.
 - In reference counting, the approximation is that an object can't be live unless it is referenced.
 - In tracing garbage collection, the approximation is that an object can't be live unless it is reachable.
- Two forms of conservatism.
 - **Temporal** conservatism occurs when garbage goes uncollected between collection cycles.
 - **Topological** conservatism occurs when different paths that share an edge in the graph of pointer relationships are not distinguished.

Run-Time Criterion for Liveness

- Defined in terms of a **root set** of variables and **reachability** from these roots.

Run-Time Criterion for Liveness

- Defined in terms of a root set of variables and reachability from these roots.
- At a garbage collection point, all globally visible variables and local variables of active methods are considered live.
 - This is a conservative notion.

Run-Time Criterion for Liveness

- Defined in terms of a root set of variables and reachability from these roots.
- At a garbage collection point, all globally visible variables and local variables of active methods are considered live.
 - This is a conservative notion.
- The root set consists of the global variables, local variables on the activation stack, and any registers used by active methods.

Run-Time Criterion for Liveness

- Defined in terms of a root set of variables and reachability from these roots.
- At a garbage collection point, all globally visible variables and local variables of active methods are considered live.
 - This is a conservative notion.
- The root set consists of the global variables, local variables on the activation stack, and any registers used by active methods.
- The set of live objects at a collection point is defined inductively as follows.
 - Heap objects directly reachable from any of the variables in the root set are live.
 - Any object directly reachable from a live object is live.
 - No other object is live.

Object Representations

- Heap objects must be self-identifying.
 - It must be easy to determine the type of an object at run-time.
 - Statically typed languages use an extra field containing type information in the hidden headers of allocated objects, and use this information to decode the format of the object.
 - Dynamically typed languages typically use tagged pointers.

Object Representations

- Heap objects must be self-identifying.
 - It must be easy to determine the type of an object at run-time.
 - Statically typed languages use an extra field containing type information in the hidden headers of allocated objects, and use this information to decode the format of the object.
 - Dynamically typed languages typically use tagged pointers.
- Actually, we can get by with less for statically typed languages.
 - All we need are the types of the root set variables.
 - From that, we can figure out the types of their referents, and decode the fields in these types, and continue this process transitively.

Object Representations

- Heap objects must be self-identifying.
 - It must be easy to determine the type of an object at run-time.
 - Statically typed languages use an extra field containing type information in the hidden headers of allocated objects, and use this information to decode the format of the object.
 - Dynamically typed languages typically use tagged pointers.
- Actually, we can get by with less for statically typed languages.
 - All we need are the types of the root set variables.
 - From that, we can figure out the types of their referents, and decode the fields in these types, and continue this process transitively.
- A class of **conservative garbage collectors** can work even without this minimal information, at the cost of a more conservative approximation to true garbage.