

ELF Relocatable Object File Format

ELF header (16 B)		Bootstrapping information for file
<code>.text</code>		Machine code of compiled module
<code>.rodata</code>		Read-only data (e.g., <code>printf</code> format strings, jump tables)
<code>.data</code>		Initialized global / static variables
<code>.bss</code>		Uninitialized static variables + those initialized to 0
<code>.symtab</code>		Symbol table
<code>.rel.text</code>		List of <code>.text</code> locations that need to be modified
<code>.rel.data</code>		List of <code>.data</code> locations that need to be modified
<code>.debug</code>	optional	Debugging symbol table
<code>.line</code>	optional	Mapping between source line #s and <code>.text</code> instructions
<code>.strtab</code>		String table for symbols in <code>.symtab</code> , <code>.debug</code> , and section names
Section Header Table		Fixed-size entries describing each section

The Tasks of The Linker

- Input: A collection $M = \{M_1, \dots, M_k\}$ of relocatable object modules.
 - Let module $M_i = (T_i, D_i, \Sigma_i)$, with T_i denoting its `.text` section, D_i its `.data` section, and Σ_i its symbol table.
 - Σ_i contains information about the symbols defined in or referenced by module M_i .

The Tasks of The Linker

- Input: A collection $M = \{M_1, \dots, M_k\}$ of relocatable object modules.
 - Let module $M_i = (T_i, D_i, \Sigma_i)$, with T_i denoting its `.text` section, D_i its `.data` section, and Σ_i its symbol table.
 - Σ_i contains information about the symbols defined in or referenced by module M_i .
- Symbol resolution: Match each reference to a **link-time symbol** in each module with exactly one symbol definition from the set of defined symbols across all the modules in M .

The Tasks of The Linker

- Input: A collection $M = \{M_1, \dots, M_k\}$ of relocatable object modules.
 - Let module $M_i = (T_i, D_i, \Sigma_i)$, with T_i denoting its `.text` section, D_i its `.data` section, and Σ_i its symbol table.
 - Σ_i contains information about the symbols defined in or referenced by module M_i .
- Symbol resolution: Match each reference to a link-time symbol in each module with exactly one symbol definition from the set of defined symbols across all the modules in M .
- Relocation: Merge `.text` and `.data` sections from multiple relocatable object modules in a coordinated manner.
 - Create a single `.text` section T from $\{T_1, \dots, T_k\}$, and a single `.data` section D from $\{D_1, \dots, D_k\}$.
 - Adjust any text and data symbol references that need to change as a result of this merge.

The Tasks of The Linker

- Input: A collection $M = \{M_1, \dots, M_k\}$ of relocatable object modules.
 - Let module $M_i = (T_i, D_i, \Sigma_i)$, with T_i denoting its `.text` section, D_i its `.data` section, and Σ_i its symbol table.
 - Σ_i contains information about the symbols defined in or referenced by module M_i .
- Symbol resolution: Match each reference to a link-time symbol in each module with exactly one symbol definition from the set of defined symbols across all the modules in M .
- Relocation: Merge `.text` and `.data` sections from multiple relocatable object modules in a coordinated manner.
 - Create a single `.text` section T from $\{T_1, \dots, T_k\}$, and a single `.data` section D from $\{D_1, \dots, D_k\}$.
 - Adjust any text and data symbol references that need to change as a result of this merge.
- Output: The binary executable or shared object file.

Link-Time Symbols

- A **link-time symbol** in a relocatable object module is a procedure name, global variable, or `static` variable that is declared, defined, or referenced in the module.

Link-Time Symbols

- A link-time symbol in a relocatable object module is a procedure name, global variable, or `static` variable that is declared, defined, or referenced in the module.
- Method-local non-`static` variables are of no interest to the linker.
 - These are compile-time symbols. They are not even described in the module's symbol table.
 - The compiler binds such symbols to an offset in the activation record for the method in which the symbol is defined.
 - An activation record is instantiated on the run-time stack for each invocation of the method.

Link-Time Symbols

- A link-time symbol in a relocatable object module is a procedure name, global variable, or `static` variable that is declared, defined, or referenced in the module.
- Method-local non-`static` variables are of no interest to the linker.
 - These are compile-time symbols. They are not even described in the module's symbol table.
 - The compiler binds such symbols to an offset in the activation record for the method in which the symbol is defined.
 - An activation record is instantiated on the run-time stack for each invocation of the method.
- Procedure-local `static` variables are allocated in the `.data` or `.bss` sections of the module.
 - These are also compile-time symbols, but they are described in the module's symbol table as they require relocation.
 - The compiler ensures that there is only one definition per module for each procedure-local static variable, and creates a unique name for it.

More on Link-Time Symbols

- Useful Boolean function on symbols:
 - $\text{DEF}(s, m) = \text{true}$, if module m contains a definition of symbol s ; **false**, otherwise.
 - $\text{USE}(s, m) = \text{true}$, if module m contains one or more references to symbol s ; **false**, otherwise.
 - $\text{VISIBLE}(s) = \text{true}$, if s is a non-static symbol; **false**, otherwise.

More on Link-Time Symbols

- Useful Boolean function on symbols:
 - $\text{DEF}(s, m) = \text{true}$, if module m contains a definition of symbol s ; **false**, otherwise.
 - $\text{USE}(s, m) = \text{true}$, if module m contains one or more references to symbol s ; **false**, otherwise.
 - $\text{VISIBLE}(s) = \text{true}$, if s is a non-static symbol; **false**, otherwise.
- With reference to module m , link-time symbol s is said to be:
 - **Local**, if $\text{DEF}(s, m) \wedge \neg \text{VISIBLE}(s)$.
 - **Global**, if $\text{DEF}(s, m) \wedge \text{VISIBLE}(s)$.
 - **External**, if $\text{USE}(s, m) \wedge \neg \text{DEF}(s, m)$.

More on Link-Time Symbols

- Useful Boolean function on symbols:
 - $\text{DEF}(s, m) = \text{true}$, if module m contains a definition of symbol s ; **false**, otherwise.
 - $\text{USE}(s, m) = \text{true}$, if module m contains one or more references to symbol s ; **false**, otherwise.
 - $\text{VISIBLE}(s) = \text{true}$, if s is a non-static symbol; **false**, otherwise.
- With reference to module m , link-time symbol s is said to be:
 - Local, if $\text{DEF}(s, m) \wedge \neg \text{VISIBLE}(s)$.
 - Global, if $\text{DEF}(s, m) \wedge \text{VISIBLE}(s)$.
 - External, if $\text{USE}(s, m) \wedge \neg \text{DEF}(s, m)$.
- The linker resolves external symbol references with global symbol definitions across modules.
 - When the compiler encounters an external symbol, it generates a link-time symbol in the module's symbol table, and passes it on to the linker to resolve.
 - The linker tries to find a *unique* definition for this symbol reference among the symbol definitions across all the modules that it is linking.

Link-Time Symbols: Examples

main.c

```
int sum(int*, int);

int array[2] = {1,2};

int main(void) {
    int val = sum(array, 2);
    return val;
}
```

sum.c

```
int sum(int *a, int n) {
    int s = 0;
    for (int i = 0; i < n; i++)
        s += a[i];
    return s;
}
```

- The symbol `sum` is **external**.
- The symbol `array` is **global**.
- The symbol `main` is **global**.

- The symbol `sum` is **global**.

Link-Time Symbols: Examples

main.c

```
int sum(int*, int);

int array[2] = {1,2};

int main(void) {
    int val = sum(array, 2);
    return val;
}
```

sum.c

```
int sum(int *a, int n) {
    int s = 0;
    for (int i = 0; i < n; i++)
        s += a[i];
    return s;
}
```

- The symbol sum is **global**.

- The symbol sum is **external**.
- The symbol array is **global**.
- The symbol main is **global**.

sum1.c

```
static int sum(int *a, int n) {
    int s = 0;
    for (int i = 0; i < n; i++)
        s += a[i];
    return s;
}
```

- The symbol sum is **local**.

Rules for Unique Resolution

- The *global* symbols of module m are further divided into *strong* and *weak* symbols.
 - Function names and initialized variables are **strong** symbols.
 - Uninitialized variables are **weak** symbols.

Rules for Unique Resolution

- The *global* symbols of module m are further divided into *strong* and *weak* symbols.
 - Function names and initialized variables are strong symbols.
 - Uninitialized variables are weak symbols.
- Let s and w be the total number of strong and weak definitions matching an external reference r .
 - If $s = w = 0$, declare a **link-time error**: no definition found.
 - If $s > 1$, declare a **link-time error**: multiple strong definitions.
 - If $s = 1$, resolve r to the **unique** strong definition.
 - If $s = 0 \wedge w = 1$, resolve r to the **unique** weak definition.
 - If $s = 0 \wedge w > 1$, resolve r to an **arbitrary** weak definition.