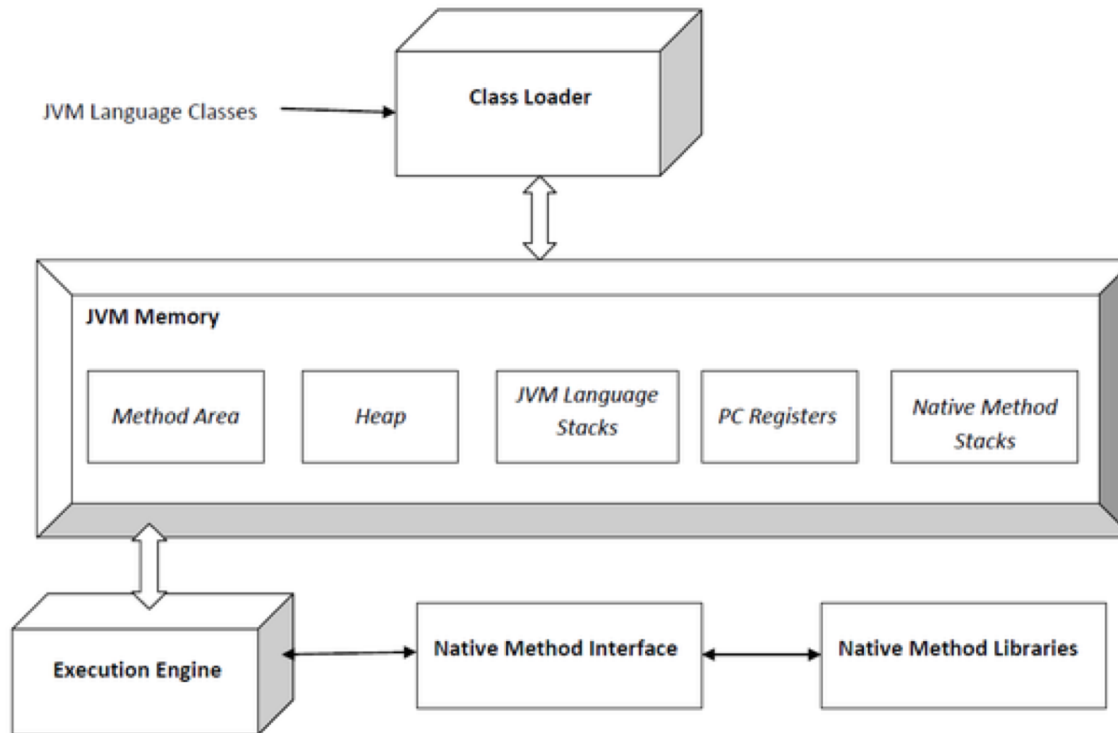


# Architecture of the Java Virtual Machine



# Linking In The Java Virtual Machine

- Java's approach to linking is completely dynamic.
- Individual `.class` files are loaded, linked, and initialized when they are first referenced.
  - Allows for greater flexibility and portability for the user.
  - Makes the necessary run-time data structures and algorithms more elaborate.



# Linking In The Java Virtual Machine

- Java's approach to linking is completely dynamic.
- Individual `.class` files are loaded, linked, and initialized when they are first referenced.
  - Allows for greater flexibility and portability for the user.
  - Makes the necessary run-time data structures and algorithms more elaborate.

```
% java HelloWorld "Hello!" "Hi!"
```

```
public static void HelloWorld.main(args)    args.length = 2  
                                              args[0] = "Hello!"  
                                              args[1] = "Hi!"
```

- Game plan
  - The `class` file format.
  - JVM startup and exit actions.
  - Class loading, linking, initialization (simplified: fast path only).

# The `class` File Format

- [Ref: JVM SE16, Chapter 4]
- A `class` file consists of a stream of 8-bit bytes.
- 16-bit and 32-bit quantities are constructed from two or four 8-bit bytes.
- Multibyte data items are always stored in big-endian order.
- Shorthand
  - `u1` for unsigned one-byte quantity.
  - `u2` for unsigned two-byte quantity.
  - `u4` for unsigned four-byte quantity.
- A class file consists of a single `ClassFile` structure.
  - Essentially, a sequence of self-describing nested tables.

# The class File Format

```
ClassFile {
    u4 magic;                // 0xCAFEBAFE
    u2 minor_version;        // Minor version number, in {0, 65535}
    u2 major_version;        // Major version number, in [45,60]

    // String constants, class and interface names, field names, etc.
    u2 constant_pool_count;
    cp_info constant_pool[constant_pool_count-1];

    u2 access_flags;
    u2 this_class;           // index into constant_pool
    u2 super_class;          // index into constant_pool

    u2 interfaces_count;
    u2 interfaces[interfaces_count]; // indices into constant_pool

    // Complete description of a field
    u2 fields_count;
    field_info fields[fields_count];

    // Complete description of a method
    u2 methods_count;
    method_info methods[methods_count];

    // Information used both for JVM function and for tools
    u2 attributes_count;
    attribute_info attributes[attributes_count];
}
```

# Constants, Fields, and Methods

- Constant pool
  - The central repository of names found in the class.
  - Consists of a number of variable-sized (but self-describing) entries of type `cp_info`, describing 17 kinds of possible constants.

```
cp_info {  
    u1 tag;  
    u1 info[];  
}
```

# Constants, Fields, and Methods

- Constant pool
  - The central repository of names found in the class.
  - Consists of a number of variable-sized (but self-describing) entries of type `cp_info`, describing 17 kinds of possible constants.
- Fields
  - Described by a `field_info` structure.
  - Variable-sized, but self-describing.

```
cp_info {  
    u1 tag;  
    u1 info[];  
}
```

```
field_info {  
    u2 access_flags;  
    u2 name_index;  
    u2 descriptor_index;  
    u2 attributes_count;  
    attribute_info  
        attributes[attributes_count];  
}
```



# Constants, Fields, and Methods

- Constant pool
  - The central repository of names found in the class.
  - Consists of a number of variable-sized (but self-describing) entries of type `cp_info`, describing 17 kinds of possible constants.
- Fields
  - Described by a `field_info` structure.
  - Variable-sized, but self-describing.
- Methods
  - Described by a `method_info` structure.
  - Variable-sized, but self-describing.
  - A key attribute is `Code`, which contains the JVM code for the method.

```
cp_info {  
    u1 tag;  
    u1 info[];  
}
```

```
field_info {  
    u2 access_flags;  
    u2 name_index;  
    u2 descriptor_index;  
    u2 attributes_count;  
    attribute_info  
        attributes[attributes_count];  
}
```

```
method_info {  
    u2 access_flags;  
    u2 name_index;  
    u2 descriptor_index;  
    u2 attributes_count;  
    attribute_info  
        attributes[attributes_count];  
}
```

# Java Virtual Machine Startup and Exit

- [Ref: JVM SE16, §5.2, §5.7]
- JVM startup
  - Create an initial class (or interface) using a class loader.
  - Link this initial class.
  - Initialize this initial class.
  - Invoke the `public static` method `void main( String[ ] )` of this class.

# Java Virtual Machine Startup and Exit

- [Ref: JVM SE16, §5.2, §5.7]
- JVM startup
  - Create an initial class (or interface) using a class loader.
  - Link this initial class.
  - Initialize this initial class.
  - Invoke the `public static` method `void main(String[ ])` of this class.
- The invocation of this method drives all further execution, including the linking and creation of additional classes and interfaces, and the invocation of additional methods.

# Java Virtual Machine Startup and Exit

- [Ref: JVM SE16, §5.2, §5.7]
- JVM startup
  - Create an initial class (or interface) using a class loader.
  - Link this initial class.
  - Initialize this initial class.
  - Invoke the `public static` method `void main(String[ ])` of this class.
- The invocation of this method drives all further execution, including the linking and creation of additional classes and interfaces, and the invocation of additional methods.
- JVM exit
  - The Java Virtual Machine exits when some thread invokes the `exit` method of class `Runtime` or class `System`, or the `halt` method of class `Runtime`, and the `exit` or `halt` operation is permitted by the security manager.

# Run-Time Constant Pool

- [Ref: JVM SE16, §5.1]
- The JVM maintains a run-time constant pool for each class (or interface).
  - Plays a role analogous to that of a symbol table.
  - Constructed upon class creation from the `constant_pool` in the `class` file representation of the class.

# Run-Time Constant Pool

- [Ref: JVM SE16, §5.1]
- The JVM maintains a run-time constant pool for each class (or interface).
  - Plays a role analogous to that of a symbol table.
  - Constructed upon class creation from the `constant_pool` in the `class` file representation of the class.
- Two kinds of entries:
  - **Symbolic references**, which may later be resolved.
  - **Static constants** (i.e., string and numeric constants), which require no further processing.

# Run-Time Constant Pool

- [Ref: JVM SE16, §5.1]
- The JVM maintains a run-time constant pool for each class (or interface).
  - Plays a role analogous to that of a symbol table.
  - Constructed upon class creation from the `constant_pool` in the `class` file representation of the class.
- Two kinds of entries:
  - Symbolic references, which may later be resolved.
  - Static constants (i.e., string and numeric constants), which require no further processing.
- All static constants and some types of symbolic references are **loadable**, i.e., they may be pushed onto the evaluation stack by the `ldc` family of instructions.

# Class Creation and Loading

- [Ref: JVM SE16, §5.3]
- **Creation** of a class (or interface)  $C$  denoted by the name  $N$  consists of the construction in the method area of the JVM of an implementation-specific internal representation of  $C$ .
  - A non-array class is created by loading its binary representation using a class loader  $L$ .
  - The run-time identity of class is the pair  $\langle N, L \rangle$ .



# Class Creation and Loading

- [Ref: JVM SE16, §5.3]
- Creation of a class (or interface)  $C$  denoted by the name  $N$  consists of the construction in the method area of the JVM of an implementation-specific internal representation of  $C$ .
  - A non-array class is created by loading its binary representation using a class loader  $L$ .
  - The run-time identity of class is the pair  $\langle N, L \rangle$ .
- **Loading** a class (when  $L$  is the bootstrap class loader)
  - The JVM passes the argument  $N$  to an invocation of a method on the bootstrap class loader to search for a representation of  $C$  in a platform-dependent manner (i.e., a `class` file).
  - Then the JVM attempts to derive a class  $C$  from the `class` file.
    - Parse the `class` file into an implementation-specific internal representation.
    - If  $C$  has a direct superclass  $D$ , **resolve** the symbolic reference from  $C$  to  $D$ . Likewise for any superinterfaces.

# Class Linking

- [Ref: JVM SE16, §5.4]
- **Linking** a class (or interface) involves **verifying** and **preparing** that class, its direct superclass, and its direct superinterfaces, if necessary.
- Linking also involves **resolution** of symbolic references in the class or interface, though not necessarily at the same time as the class or interface is verified and prepared.
- The JVM specification allow an implementation to be flexible in scheduling linking activities “eagerly” or “lazily”, as long as the following invariants are maintained.
  - A class is completely loaded before it is linked.
  - A class is completely verified and prepared before it is initialized.

# Verification, Preparation, Resolution

- [Ref: JVM SE16, §5.4.1-3]
- **Verification** ensures that the binary representation of a class or interface is structurally correct.

# Verification, Preparation, Resolution

- [Ref: JVM SE16, §5.4.1-3]
- Verification ensures that the binary representation of a class or interface is structurally correct.
- **Preparation** involves creating the static fields for a class and initializing such fields to their default values.
  - This does not require the execution of any JVM code; explicit initializers for static fields are executed as part of initialization, not preparation.
  - Preparation may occur at any time following creation but must be completed prior to initialization.

# Verification, Preparation, Resolution

- [Ref: JVM SE16, §5.4.1-3]
- Verification ensures that the binary representation of a class or interface is structurally correct.
- Preparation involves creating the static fields for a class and initializing such fields to their default values.
  - This does not require the execution of any JVM code; explicit initializers for static fields are executed as part of initialization, not preparation.
  - Preparation may occur at any time following creation but must be completed prior to initialization.
- **Resolution** is the process of dynamically determining one or more concrete values from a symbolic reference in the run-time constant pool.
  - Required when a JVM instruction like `new`, `invokevirtual`, or `getfield` is executed.
  - Initially, all symbolic references in the run-time constant pool are unresolved.

# Class Initialization

- [Ref: JVM SE16, §5.5]
- **Initialization** of a class consists of executing its class initialization method (i.e., `<clinit>`).
  - Prior to initialization, a class must be linked, i.e., verified, prepared, and optionally resolved.
- Details of the initialization process are complicated.
  - Because the JVM is multithreaded, initialization of a class requires careful synchronization, since some other thread may be trying to initialize the same class at the same time.
  - There is also the possibility that initialization of a class may be requested recursively as part of the initialization of that class.
  - The implementation of the JVM is responsible for taking care of synchronization and recursive initialization following a specified procedure.