

The Dark Side of Data-Code Duality

- The stack frame organization on x86 mixes both data and pointers to code.
 - Data: Local variables.
 - Pointers to code: Return address.

Position	Contents	Frame
$8n+16(\%rbp)$	memory argument eightbyte n	Previous
	...	
$16(\%rbp)$	memory argument eightbyte 0	
$8(\%rbp)$	return address	Current
$0(\%rbp)$	previous $\%rbp$ value	
$-8(\%rbp)$	unspecified	
	...	
$0(\%rsp)$	variable size	
$-128(\%rsp)$	red zone	



The Dark Side of Data-Code Duality

- The stack frame organization on x86 mixes both data and pointers to code.
 - Data: Local variables.
 - Pointers to code: Return address.
- What happens if we write some inappropriate bit pattern into the return address location?
 - We will corrupt the stack, subvert the control flow of the program, and cause it to demonstrate unexpected behavior.

Position	Contents	Frame
$8n+16 (\%rbp)$	memory argument eightbyte n	Previous
	...	
$16 (\%rbp)$	memory argument eightbyte 0	
$8 (\%rbp)$	return address	Current
$0 (\%rbp)$	previous $\%rbp$ value	
$-8 (\%rbp)$	unspecified	
	...	
$0 (\%rsp)$	variable size	
$-128 (\%rsp)$	red zone	



The Dark Side of Data-Code Duality

- The stack frame organization on x86 mixes both data and pointers to code.
 - Data: Local variables.
 - Pointers to code: Return address.

- What happens if we write some inappropriate bit pattern into the return address location?

- We will corrupt the stack, subvert the control flow of the program, and cause it to demonstrate unexpected behavior.

Position	Contents	Frame
$8n+16 (\%rbp)$	memory argument eightbyte n	Previous
	...	
$16 (\%rbp)$	memory argument eightbyte 0	
$8 (\%rbp)$	return address	Current
$0 (\%rbp)$	previous $\%rbp$ value	
$-8 (\%rbp)$	unspecified	
	...	
$0 (\%rsp)$	variable size	
$-128 (\%rsp)$	red zone	

- This corruption can be unintended or malicious.
 - A common source of such corruption is known as **buffer overflow**.



Buffer Overflow Example

```
char *gets(char *s) {  
    int c;  
    char *dest = s;  
    while ((c = getchar()) != '\n' && c != EOF)  
        *dest++ = c;  
    if (c == EOF && dest == s) return NULL;  
    *dest++ = '\0';  
    return s;  
}
```

Buffer Overflow Example

```
char *gets(char *s) {  
    int c;  
    char *dest = s;  
    while ((c = getchar()) != '\n' && c != EOF)  
        *dest++ = c;  
    if (c == EOF && dest == s) return NULL;  
    *dest++ = '\0';  
    return s;  
}
```

```
void echo() {  
    char buf[8];  
    gets(buf);  
    puts(buf);  
}
```

Buffer Overflow Example


```
char *gets(char *s) {
    int c;
    char *dest = s;
    while ((c = getchar()) != '\n' && c != EOF)
        *dest++ = c;
    if (c == EOF && dest == s) return NULL;
    *dest++ = '\0';
    return s;
}
```

```
void echo() {
    char buf[8];
    gets(buf);
    puts(buf);
}
```

```
echo:
    subq    $24, %rsp        # Allocate 24 bytes on stack
    movq    %rsp, %rdi       # Compute buf as %rsp
    call    gets             # Call gets
    movq    %rsp, %rdi       # Compute buf as %rsp
    call    puts             # Call puts
    addq    $24, %rsp        # Deallocate stack space
    ret                     # Return
```

Buffer Overflow Example

```
char *gets(char *s) {  
    int c;  
    char *dest = s;  
    while ((c = getchar()) != '\n' && c != EOF)  
        *dest++ = c;  
    if (c == EOF && dest == s) return NULL;  
    *dest++ = '\0';  
    return s;  
}
```



```
void echo() {  
    char buf[8];  
    gets(buf);  
    puts(buf);  
}
```

```
echo:  
    subq    $24, %rsp        # Allocate 24 bytes on stack  
    movq    %rsp, %rdi       # Compute buf as %rsp  
    call    gets             # Call gets  
    movq    %rsp, %rdi       # Compute buf as %rsp  
    call    puts             # Call puts  
    addq    $24, %rsp        # Deallocate stack space  
    ret                    # Return
```

Consequences of Buffer Overflow

- In our example, the consequence was unexpected behavior, possibly a core dump.

Consequences of Buffer Overflow

- In our example, the consequence was unexpected behavior, possibly a core dump.
- More pernicious uses are possible.
 - Feed the program with a string that contains the byte encoding of some executable (aka the **exploit code**), plus some extra bytes that overwrite the return address with a pointer to the exploit code.
 - In this case, executing the `ret` instruction causes the program to jump to the exploit code.
 - The exploit code can now use a system call to start up a shell program.
 - Or the exploit code can perform some otherwise unauthorized task, repair the damage to the stack, and then execute a second `ret`, thereby hiding its malicious behavior from the user.

Thwarting Buffer Overflow: #1

- Stack randomization (or address-space layout randomization)
 - *Insight:*
In order to insert exploit code, the attacker needs to inject both the code and a pointer to this code as part of the attack string. We need to know the stack address of the string in order to generate the pointer value.
 - *Solution:*
Make the position of the stack vary from one run of a program to another.

```
int main() {  
    long local;  
    printf("local at %p\n", &local);  
    return 0;  
}
```
 - *Weakness:*
Still susceptible to brute force attacks.

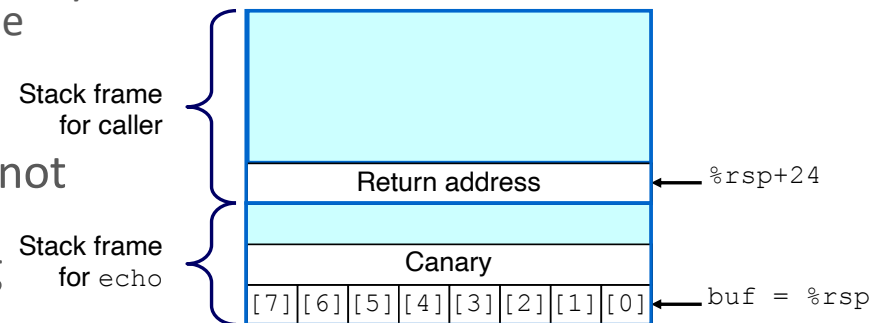
Thwarting Buffer Overflow: #2

- Stack Corruption Detection

- Insight:*
If a buffer overflows, it will leave a trace in memory.
- Solution:*
Put a stack protector into the generated code to detect buffer overflows.
 - Store a special canary value (aka guard value) in the stack frame between any local buffer and the rest of the stack state.
 - Canary value is generated randomly for each run of program.
 - Before restoring register state and returning from the function, the canary value has been checked for possible alteration.

- Weakness:*
A small performance penalty. Does not protect against other ways of corrupting the state of an executing program.

```
echo:
    subq    $24, %rsp
    movq    %fs:40, %rax
    movq    %rax, 8(%rsp)
    xorl    %eax, %eax
    movq    %rsp, %rdi
    call    gets
    movq    %rsp, %rdi
    call    puts
    movq    8(%rsp), %rax
    xorq    %fs:40, %rax
    je      .L9
    call    __stack_chk_fail
.L9:
    addq    $24, %rsp
    ret
```



Thwarting Buffer Overflow: #3

- Limiting Executable Code Regions
 - *Insight:*
The exploit won't work unless the exploit code is actually executable.
 - *Solution:*
Limit the memory regions that can hold executable code.
 - Only the `.text` section should be executable.
 - Use memory protection mechanisms in the operating system to accomplish this.
 - *Weakness:*
 - Needs to be modified for JIT compilers (or dynamic compilers in general) that generate code at run-time.
 - Still susceptible to return-oriented programming attacks.